

# INF283 | Weekly Exercise 02 | Regression

## Table of Contents:

### 1. Univariate Linear Regression

- Exercise 1.1
- Exercise 1.2
- Exercise 1.3
- Exercise 1.4
- Exercise 1.5

### 2. Multivariate Linear Regression

- Exercise 2.1
- Exercise 2.2
- Exercise 2.3

### 3. Logistic Regression

- Exercise 3.1

## 1. Univariate Linear Regression

## Linear Regression with Gradient Descent

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import pandas as pd
```

### Finding the mean squared error (MSE)

In [2]:

```
# We are now going to define a function which takes the parameters (w0 and w1) of
# a line and then finds the mean-squared error between the user-specified points
# and the line.
def compute_error_for_line_given_points(w0, w1, points):
    totalError = 0
    for i in range(0, len(points)):
        x = points[i, 0]
        y = points[i, 1]
        # accumulate 'sum of square' errors in totalError variable
        totalError += (y - (w1 * x + w0)) ** 2
    # find mean of sum of squared errors
    mse = totalError/len(points)
    return mse

# N.B.: Students who wish to do this exercise in R should implement this function in R
themselves
```

## Gradient Descent

In [3]:

```
def step_gradient(w0_current, w1_current, points, learningRate):
    #initialize the partial derivatives for the cumulative sum
    w0_par_der = 0
    w1_par_der = 0

    n = len(points)

    # computation for the summation
    for i in range(0, len(points)):
        x = points[i, 0]
        y = points[i, 1]
        # partial derivative (of MSE) with respect to w0
        w0_par_der += (y - ((w1_current * x) + w0_current))
        # partial derivative (of MSE) with respect to w1
        w1_par_der += x * (y - ((w1_current * x) + w0_current))

    # multiplication of summation results with -2/n
    w0_par_der = -(2/n) * w0_par_der
    # partial derivative (of MSE) with respect to w1
    w1_par_der = -(2/n) * w1_par_der

    # make a gradient vector from the partial derivatives
    gradient_mse = np.array([w0_par_der, w1_par_der])

    # make a vector of weights
    weight_vector = np.array([w0_current, w1_current])

    # update rule for weights
    updated_weight_vector = weight_vector - (learningRate * gradient_mse)

    # return the updated weight vector as a List
    return np.ndarray.tolist(updated_weight_vector)

# N.B. Students who wish to do this exercise in R should implement this function in R
themselves
```

## Running Gradient Descent Iteratively

In [4]:

```
def gradient_descent_runner(points, starting_w0, starting_w1, learning_rate, num_iterations):
    w0 = starting_w0
    w1 = starting_w1
    for i in range(num_iterations):
        w0, w1 = step_gradient(w0, w1, points, learning_rate)
        mse = compute_error_for_line_given_points(w0, w1, points)
        print(f'Iteration {i+1}: w0={w0:0.5f}, w1={w1:0.5f}, mse={mse:0.5f}')
    return [w0, w1, mse]
# N.B.: Students who wish to do this exercise in R should implement this function in R themselves
```

## Bringing it all together

In [5]:

```
np.random.seed(2)


# generate 100 x values from 0 to 2 randomly, then sort them in ascending order
X = 2 * np.random.rand(100, 1)
X.sort(axis=0)

# generate y values and add noise to it
y = 4 + 3 * X + np.random.rand(100, 1)

# let us plot the data
plt.scatter(X, y)
```

Out[5]:

<matplotlib.collections.PathCollection at 0x269541d7fd0>



In [8]:

```
# combine the x and y values into a single array called points
points = np.column_stack((X, y))

num_iterations = 100
learning_rate = 0.1
initial_w0 = 0 # initial y-intercept guess
initial_w1 = 0 # initial slope guess
early_stop = 0.00064
[w0, w1, mse] = gradient_descent_runner_early_stop(points, initial_w0, initial_w1, learning_rate, num_iterations, early_stop)
```

Iteration 1:  $w_0=1.44691$ ,  $w_1=1.50194$ ,  $mse=20.14308$   
Iteration 2:  $w_0=2.32806$ ,  $w_1=2.39849$ ,  $mse=7.47831$   
Iteration 3:  $w_0=2.86801$ ,  $w_1=2.93046$ ,  $mse=2.87255$   
Iteration 4:  $w_0=3.20208$ ,  $w_1=3.24295$ ,  $mse=1.19394$   
Iteration 5:  $w_0=3.41183$ ,  $w_1=3.42341$ ,  $mse=0.57873$   
Iteration 6:  $w_0=3.54643$ ,  $w_1=3.52453$ ,  $mse=0.35003$   
Iteration 7:  $w_0=3.63550$ ,  $w_1=3.57805$ ,  $mse=0.26199$   
Iteration 8:  $w_0=3.69691$ ,  $w_1=3.60309$ ,  $mse=0.22532$   
Iteration 9:  $w_0=3.74143$ ,  $w_1=3.61118$ ,  $mse=0.20757$   
Iteration 10:  $w_0=3.77555$ ,  $w_1=3.60925$ ,  $mse=0.19693$   
Iteration 11:  $w_0=3.80321$ ,  $w_1=3.60148$ ,  $mse=0.18912$   
Iteration 12:  $w_0=3.82676$ ,  $w_1=3.59037$ ,  $mse=0.18255$   
Iteration 13:  $w_0=3.84765$ ,  $w_1=3.57743$ ,  $mse=0.17664$   
Iteration 14:  $w_0=3.86674$ ,  $w_1=3.56358$ ,  $mse=0.17117$   
Iteration 15:  $w_0=3.88456$ ,  $w_1=3.54934$ ,  $mse=0.16604$   
Iteration 16:  $w_0=3.90144$ ,  $w_1=3.53503$ ,  $mse=0.16122$   
Iteration 17:  $w_0=3.91758$ ,  $w_1=3.52085$ ,  $mse=0.15667$   
Iteration 18:  $w_0=3.93309$ ,  $w_1=3.50690$ ,  $mse=0.15238$   
Iteration 19:  $w_0=3.94807$ ,  $w_1=3.49325$ ,  $mse=0.14833$   
Iteration 20:  $w_0=3.96257$ ,  $w_1=3.47992$ ,  $mse=0.14451$   
Iteration 21:  $w_0=3.97662$ ,  $w_1=3.46694$ ,  $mse=0.14090$   
Iteration 22:  $w_0=3.99025$ ,  $w_1=3.45430$ ,  $mse=0.13750$   
Iteration 23:  $w_0=4.00348$ ,  $w_1=3.44201$ ,  $mse=0.13428$   
Iteration 24:  $w_0=4.01632$ ,  $w_1=3.43005$ ,  $mse=0.13125$   
Iteration 25:  $w_0=4.02880$ ,  $w_1=3.41844$ ,  $mse=0.12838$   
Iteration 26:  $w_0=4.04092$ ,  $w_1=3.40715$ ,  $mse=0.12568$   
Iteration 27:  $w_0=4.05269$ ,  $w_1=3.39618$ ,  $mse=0.12313$   
Iteration 28:  $w_0=4.06412$ ,  $w_1=3.38552$ ,  $mse=0.12072$   
Iteration 29:  $w_0=4.07523$ ,  $w_1=3.37516$ ,  $mse=0.11844$   
Iteration 30:  $w_0=4.08603$ ,  $w_1=3.36510$ ,  $mse=0.11630$   
Iteration 31:  $w_0=4.09651$ ,  $w_1=3.35533$ ,  $mse=0.11427$   
Iteration 32:  $w_0=4.10670$ ,  $w_1=3.34583$ ,  $mse=0.11236$   
Iteration 33:  $w_0=4.11660$ ,  $w_1=3.33660$ ,  $mse=0.11055$   
Iteration 34:  $w_0=4.12622$ ,  $w_1=3.32763$ ,  $mse=0.10885$   
Iteration 35:  $w_0=4.13556$ ,  $w_1=3.31892$ ,  $mse=0.10724$   
Iteration 36:  $w_0=4.14464$ ,  $w_1=3.31046$ ,  $mse=0.10572$   
Iteration 37:  $w_0=4.15346$ ,  $w_1=3.30223$ ,  $mse=0.10429$   
Iteration 38:  $w_0=4.16203$ ,  $w_1=3.29424$ ,  $mse=0.10293$   
Iteration 39:  $w_0=4.17035$ ,  $w_1=3.28648$ ,  $mse=0.10166$   
Iteration 40:  $w_0=4.17844$ ,  $w_1=3.27894$ ,  $mse=0.10045$   
Iteration 41:  $w_0=4.18630$ ,  $w_1=3.27162$ ,  $mse=0.09931$   
Iteration 42:  $w_0=4.19393$ ,  $w_1=3.26450$ ,  $mse=0.09824$   
Iteration 43:  $w_0=4.20135$ ,  $w_1=3.25758$ ,  $mse=0.09722$   
Iteration 44:  $w_0=4.20856$ ,  $w_1=3.25086$ ,  $mse=0.09627$   
Iteration 45:  $w_0=4.21556$ ,  $w_1=3.24433$ ,  $mse=0.09536$   
Iteration 46:  $w_0=4.22236$ ,  $w_1=3.23799$ ,  $mse=0.09451$   
Iteration 47:  $w_0=4.22897$ ,  $w_1=3.23183$ ,  $mse=0.09371$   
Iteration 48:  $w_0=4.23539$ ,  $w_1=3.22584$ ,  $mse=0.09295$   
Iteration 49:  $w_0=4.24163$ ,  $w_1=3.22002$ ,  $mse=0.09223$   
Iteration 50:  $w_0=4.24769$ ,  $w_1=3.21437$ ,  $mse=0.09155$

## Exercise 1.1

In [7]:

```
# TODO:
# Paste your code below for the modified gradient_descent_runner_early_stop function
def gradient_descent_runner_early_stop(points, starting_w0, starting_w1, learning_rate,
    num_iterations, early_stop):
    w0 = starting_w0
    w1 = starting_w1
    mse_t0 = 1e4 # Set to a high number for case 0
    for i in range(num_iterations):
        w0, w1 = step_gradient(w0, w1, points, learning_rate)
        mse_t1 = compute_error_for_line_given_points(w0, w1, points)
        if(mse_t0 - mse_t1 < early_stop):
            break;
        else:
            print(f'Iteration {i+1}: w0={w0:0.5f}, w1={w1:0.5f}, mse={mse_t1:0.5f}')
            mse_t0 = mse_t1
    return [w0, w1, mse_t0]
# N.B.: Students who wish to do this exercise in R should implement this function in R
themselves
```

## Exercise 1.2

In the program above we had set the learning rate to 0.1. Using the original `gradient_descent_runner` function, first set the number of iterations to 100. Then try to run the code with two different values of learning rates:

1. a learning rate of 0.001
2. a learning rate of 1

Explain what you observe.

In [9]:

```
%%capture
num_iterations = 100
learning_rate = [1, 1e-3]
initial_w0 = 0 # initial y-intercept guess
initial_w1 = 0 # initial slope guess
for lr in learning_rate:
    print("-----" + "\n" + "Learning rate: " + str(lr))
    [w0, w1, mse] = gradient_descent_runner(points, initial_w0, initial_w1, lr, num_ite
rations)
```

### TODO

When learning rate is set to 1 the mse (drastically) increases. </br> When learning rate is set to 0.001 the mse decreases, however it decreases slowly. </br> Conclusion: Learning rate = 1 is too high, while learning rate = 0.001 is too low.

## Exercise 1.3

In [10]:

```
# TODO
# Write your solution here
shape = X.shape
obj_type = X.dtype
ones = np.ones((shape), obj_type)
X_ones = np.hstack((X, ones))
```

In [11]:

```
X_trans = X_ones.T
```

In [12]:

```
X_trans_times_X = X_trans.dot(X_ones)
```

In [13]:

```
from numpy.linalg import inv
```

In [14]:

```
X_T_times_X_inv = inv(X_trans_times_X)
```

In [15]:

```
w_hat = ((X_T_times_X_inv.dot(X_trans)).dot(y))
```

In [16]:

```
w_hat
```

Out[16]:

```
array([[3.02129039],
       [4.45478709]])
```

It seems like the values are more or less the same, so both methods obviously work.

## Exercise 1.4

To brush up your calculus skills, derive the partial derivate of MSE that has  $L2$  penalty term included in it. In other words, we want you compute the following partial derivatives:

$$\frac{\partial}{\partial w_0} \left( \frac{1}{n} \sum_{i=1}^n [y_i - (w_1 x_i + w_0)]^2 + \lambda w_0^2 \right)$$

$$\frac{\partial}{\partial w_1} \left( \frac{1}{n} \sum_{i=1}^n [y_i - (w_1 x_i + w_0)]^2 + \lambda w_1^2 \right)$$

### What to submit ¶

A derivation of both the gradient equations.

$$\frac{\partial}{\partial w_0} \left( \frac{1}{n} \sum_{i=1}^n [y_i - (w_1 x_i + w_0)]^2 + \lambda \cdot w_0^2 \right) = -\frac{2}{n} \left( \sum_{i=1}^n y_i - (w_1 x_i + w_0) + 2\lambda w_0 \right)$$

$$\frac{\partial}{\partial w_1} \left( \frac{1}{n} \sum_{i=1}^n [y_i - (w_1 x_i + w_0)]^2 + \lambda \cdot w_1^2 \right) = -\frac{2}{n} \left( \sum_{i=1}^n y_i x_i - (w_1 x_i + w_0) x_i + 2\lambda w_1 \right)$$

## Linear Regression with *sklearn* Machine Learning Library

In [17]:

```
# import LinearRegression class from sklearn.linear_model module
from sklearn.linear_model import LinearRegression

# make a lin_reg object form the LinearRegression class
lin_reg = LinearRegression()

# use the fit method of LinearRegression class to fit a straight line through the data
lin_reg.fit(X, y)
```

Out[17]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

In [18]:

```
print('slope w1:', lin_reg.coef_)
print('y-intercept w0:', lin_reg.intercept_)
```

```
slope w1: [[3.02129039]]
y-intercept w0: [4.45478709]
```



In [19]:

```
# plot the original data points as a scatter plot
plt.scatter(X, y, label='original data')


# plot the line that fits these points. Use the values of m and b as provided by the fit method
y_ = lin_reg.coef_*X + lin_reg.intercept_

# you can also get y_ by using the predict method. Uncomment the line below:
#y_ = lin_reg.predict(X)

plt.plot(X, y_, color='r', label='predicted fit')
plt.xlabel('x'); plt.ylabel('y')
plt.legend(loc='best')
```

Out[19]:

<matplotlib.legend.Legend at 0x26955ba9630>



## Exercise 1.5

1. The *fit* method from the **sklearn library** seems to give more or less the same values for *b* and *m* as "my" implementation of the **Normal Equation**

In [20]:

```
# Answer to 2.
test_data_point = 3
pred_value = float(lin_reg.predict(test_data_point))
actual_value = 4 + 3*test_data_point
print("Predicted value: " + str(pred_value) + '\n' + "Actual value: " + str(actual_value))
```

Predicted value: 13.518658256328946

Actual value: 13

Pretty fair prediction

## 2. Multivariate Linear Regression

In [21]:

```
# make a dataframe of the data
df = pd.read_csv('movies.csv')

# show first five rows of df
df.head(n=5)
```

Out[21]:

	revenue	production_cost	promotional_cost	book_sales
0	85.099998	8.5	5.100000	4.7
1	106.300003	12.9	5.800000	8.8
2	50.200001	5.2	2.100000	15.1
3	130.600006	10.7	8.399999	12.2
4	54.799999	3.1	2.900000	10.6

In [22]:

```
# Extract the first column and set it to the output or dependent variable y
y = df[['revenue']]

# Remove the first column and set the rest of the dataframe to X. This is the set of independent variables
X = df.drop(columns=['revenue'])

# show first five rows of X
X.head(n=5)
```

Out[22]:

	production_cost	promotional_cost	book_sales
0	8.5	5.100000	4.7
1	12.9	5.800000	8.8
2	5.2	2.100000	15.1
3	10.7	8.399999	12.2
4	3.1	2.900000	10.6

In [23]:

```
# show first five rows of y
y.head(n=5)
```

Out[23]:

	revenue
0	85.099998
1	106.300003
2	50.200001
3	130.600006
4	54.799999

In [24]:

```
from sklearn.linear_model import LinearRegression

# make a lin_reg object form the LinearRegression class
lin_reg = LinearRegression()

# use the fit method of LinearRegression class to fit a straight line through the data
lin_reg.fit(X, y)

# Display the Learned parameters
lin_reg.intercept_, lin_reg.coef_
```

Out[24]:

```
(array([7.67602854]), array([[3.66160401, 7.62105126, 0.82846807]]))
```

## Exercise 2.1

In [25]:

```
## TODO
## Write your code here
prod_c = 23
prom_c = 12
book_s = 10
X_test = np.array([prod_c, prom_c, book_s]).reshape(1, -1)
lin_reg.predict(X_test)
```

Out[25]:

```
array([[191.6302165]])
```

## Multivariate Regression with Polynomial basis

In [26]:

```
# define the number of points to generate as k
k = 100

# define a seed value. It is important to define the seed value
# so that the random numbers generated are the same every time
# this code is executed.
np.random.seed(10)

# generate k x-axis values from -3 to +3
X = 6 * np.random.rand(k, 1) - 3

# sort the numbers in ascending order. This helps when we are plotting the data.
# Without this line, your plots will be all jumbled up
X.sort(axis=0)

# generate k y-axis values
y = 0.5 * X**2 + X + 2 + np.random.rand(k, 1)
```


Let us now plot the data:

In [27]:

```
plt.scatter(X, y)
plt.xlabel('x'); plt.ylabel('y')
```

Out[27]:

Text(0,0.5,'y')



In [28]:

```
from sklearn.preprocessing import PolynomialFeatures

poly_features = PolynomialFeatures(degree=15, include_bias=False)

# generate polynomial features upto degree 2 from the vector X
X_poly = poly_features.fit_transform(X)
```

In [29]:

```
# display 4 original data points
X[1:5]
```

Out[29]:

```
array([[ -2.8754883 ],
       [ -2.84760131],
       [ -2.7643094 ],
       [ -2.76024475]])
```

In [30]:

```
# Display the transformed data.
# You will now see the original X data alongside its corresponding 2nd-degree polynomial feature
X_poly[1:5]
```

Out[30]:

```
array([[ -2.87548830e+00,  8.26843299e+00, -2.37757823e+01,
         6.83669840e+01, -1.96588463e+02,  5.65287826e+02,
        -1.62547853e+03,  4.67404451e+03, -1.34401603e+04,
         3.86470238e+04, -1.11129065e+05,  3.19550326e+05,
        -9.18863225e+05,  2.64218046e+06, -7.59755900e+06],
       [ -2.84760131e+00,  8.10883321e+00, -2.30907240e+01,
         6.57531760e+01, -1.87238830e+02,  5.33181537e+02,
        -1.51828844e+03,  4.32348015e+03, -1.23115477e+04,
         3.50583794e+04, -9.98322871e+04,  2.84282551e+05,
        -8.09523365e+05,  2.30519979e+06, -6.56428994e+06],
       [ -2.76430940e+00,  7.64140644e+00, -2.11232116e+01,
         5.83910924e+01, -1.61411045e+02,  4.46190069e+02,
        -1.23340740e+03,  3.40951967e+03, -9.42496726e+03,
         2.60535256e+04, -7.20200055e+04,  1.99085578e+05,
        -5.50334134e+05,  1.52129382e+06, -4.20532680e+06],
       [ -2.76024475e+00,  7.61895107e+00, -2.10301697e+01,
         5.80484154e+01, -1.60227834e+02,  4.42268036e+02,
        -1.22076802e+03,  3.36961853e+03, -9.30097184e+03,
         2.56729587e+04, -7.08636494e+04,  1.95601016e+05,
        -5.39906677e+05,  1.49027457e+06, -4.11352255e+06]])
```

In [31]:

```
lin_reg = LinearRegression()

# Now we fit a linear model to the X_poly (the transformed features set) and y
lin_reg.fit(X_poly, y)

# show the values of intercept and learned co-efficients
lin_reg.intercept_, lin_reg.coef_
```

Out[31]:

```
(array([2.43771889]),
 array([[ 4.70240062e-01,  9.81483416e-01,  1.45901724e+00,
        -5.43920910e-01, -1.30913386e+00,  2.30186856e-01,
         5.62777734e-01, -4.49943278e-02, -1.28760085e-01,
         3.53201168e-03,  1.58848833e-02,  1.15315775e-05,
        -9.84782281e-04, -1.01041607e-05,  2.36325364e-05]]))
```

In [32]:


```
y_ = lin_reg.predict(X_poly)

plt.scatter(X, y, label='original data')
plt.plot(X, y_, color='r', label='predicted fit')
plt.legend(loc='best')

plt.xlabel('x')
plt.ylabel('y')
```

Out[32]:

Text(0,0.5,'y')



## Regularization with Ridge Penalty

In [33]:

```
from sklearn.linear_model import Ridge
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures

# define the number of points to generate
k = 100
np.random.seed(10)

# generate k x-axis values from -3 to +3
X = 6 * np.random.rand(k, 1) - 3
X.sort(axis=0)

# generate k y-axis values
y = 0.5 * X**2 + X + 2 + np.random.randn(k, 1)

# Create polynomial feature (degree 15)
poly_features = PolynomialFeatures(degree=15, include_bias=False)
X_poly = poly_features.fit_transform(X)

# Create Ridge regression object from Ridge class
ridge_reg = Ridge(alpha=5e-2)
# Fit data using Ridge regression
ridge_reg.fit(X_poly, y)

# Create Linear regression object from LinearRegress class (this is just for comparison)
lin_reg = LinearRegression()
# Fit data using Linear regression
lin_reg.fit(X_poly, y)

y_predict_ridge = ridge_reg.predict(X_poly)
y_predict_linear = lin_reg.predict(X_poly)

plt.scatter(X, y, label='data')
plt.plot(X, y_predict_linear, color='b', label='Linear regression fit')
plt.plot(X, y_predict_ridge, color='r', label='Ridge regression fit')


plt.xlabel('x')
plt.ylabel('y')

plt.legend(loc='best')
```

```
C:\Users\Sindr\Anaconda3\lib\site-packages\sklearn\linear_model\ridge.py:12: LinAlgWarning: scipy.linalg.solve
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number9.211483e-17
  overwrite_a=True).T
```

Out[33]:

<matplotlib.legend.Legend at 0x26955f7d898>



## Exercise 2.2

*TODO*

Observing that when changing(/decreasing) the **regularization parameter** the model starts overfitting

## Linear Regression with Radial Basis Functions



In [34]:

```
# Set random seed
np.random.seed(0)
m = 100


# Create random set of m x values between -6 and +6
X = np.random.rand(m, 1)*12 - 6
X.sort(axis =0)

# Create a non-linear dataset with random noise
y = 0.5*np.cos(X) + np.sin(X) + 4*np.cos(2*X) + np.exp(np.cos(3*X)) + 3*np.random.rand(m,1)

# plot it
plt.scatter(X, y)
```

Out[34]:

<matplotlib.collections.PathCollection at 0x26955fe5fd0>



In [35]:

```
from sklearn.metrics.pairwise import rbf_kernel
from sklearn.linear_model import LinearRegression

# find the transformation of X using Radial Basis Functions
# Each point in X is now modeled as vector of 100 values.
# See the X_RBF.shape and X_RBF to find how rbf_kernel transformed
# the original datapoints
X_RBF = rbf_kernel(X, X, gamma=0.1)

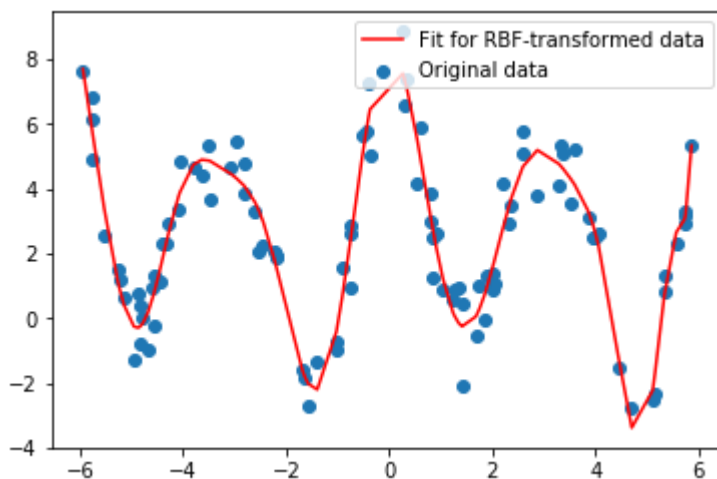
# Fit a linear regression model to the RBF-transformed data
clf = LinearRegression()
clf.fit(X_RBF, y)

# find the predicted values
y_ = clf.predict(X_RBF)

# plot original data and predicted fit
plt.scatter(X, y, label='Original data')
plt.plot(X, y_, color='r', label='Fit for RBF-transformed data')
plt.legend(loc='best')
```

Out[35]:

<matplotlib.legend.Legend at 0x26958529048>




## Exercise 2.3

In [36]:

```
# TODO
# Paste your solution here
X_RBF_0th = X_RBF[0]
X_RBF_49th = X_RBF[49]
X_RBF_99th = X_RBF[99]
ans = np.array([X_RBF_0th, X_RBF_49th, X_RBF_99th])

plt.plot(X, ans.T)
plt.show()
```



### 3. Logistic Regression

In [37]:

```
from sklearn import datasets
iris = datasets.load_iris()

# iris is a dictionary of key-value pairs. Each key-value pairs contains some informati
on about the dataset.
# Lets display a list of these keys and see what they hold
list(iris.keys())
```

Out[37]:

```
['data', 'target', 'target_names', 'DESCR', 'feature_names']
```

In [38]:

```
# Let us get the petal width. It is present in the 4th column of data
X = iris["data"][:, 3:]
X.sort(axis=0)


# Lets define a binaray variable that encodes whether a flower is Iris-Virginica or not
# Iris_virginca flower is encoded as a 2 in target
y = (iris["target"] == 2).astype(np.int) # 1 if Iris-Virginica, else 0
```

In [39]:

```
plt.scatter(X, y)
plt.xlabel('Petal width (cm)')
plt.ylabel('Iris-Virginica(1) \n Not Iris-Virginica(0)')
```

Out[39]:

Text(0,0.5,'Iris-Virginica(1) \n Not Iris-Virginica(0)')



In [40]:

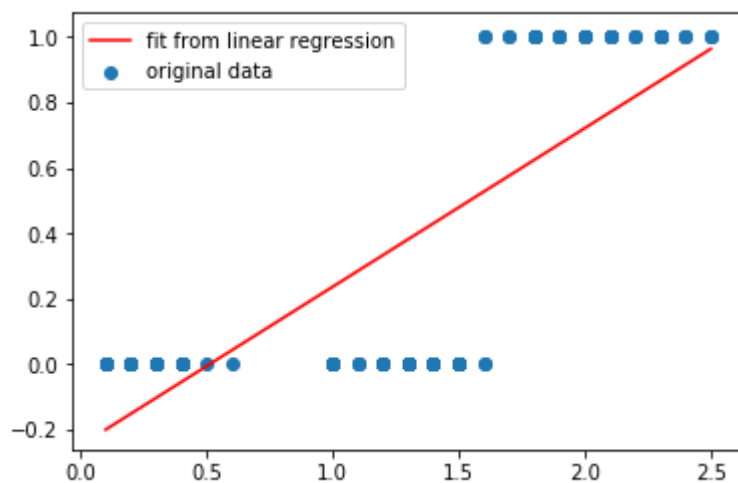
```
from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()
lin_reg.fit(X, y)
y_ = lin_reg.predict(X)

plt.scatter(X, y, label='original data')
plt.plot(X, y_, color='r', label='fit from linear regression')
plt.legend(loc='best')
```

Out[40]:

<matplotlib.legend.Legend at 0x26958593c18>



In [41]:

```
from sklearn.linear_model import LogisticRegression
log_reg = LogisticRegression()
log_reg.fit(X, y)
```

Out[41]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False)
```

In [42]:

```
# we generate X_new which is vector of closely spaced points form 0 to 3
# This vector will help us plot the model
X_new = np.linspace(0, 3, 1000).reshape(-1, 1)

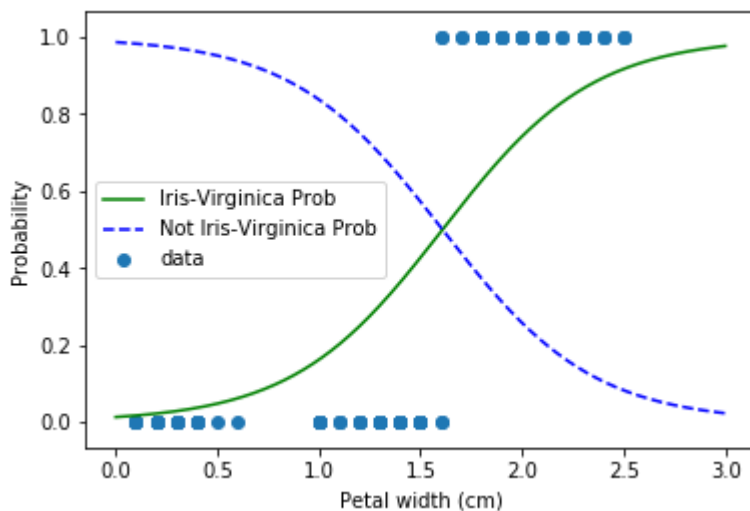
# make a vector of prediction probability values for all datapoints in X_new
y_proba = log_reg.predict_proba(X_new)

plt.plot(X_new, y_proba[:, 1], "g-", label="Iris-Virginica Prob")
plt.plot(X_new, y_proba[:, 0], "b--", label="Not Iris-Virginica Prob")
plt.scatter(X, y, label='data')

plt.xlabel('Petal width (cm)')
plt.ylabel('Probability')
plt.legend(loc='best')
```

Out[42]:

<matplotlib.legend.Legend at 0x26958614e48>



In [43]:

```
log_reg.predict([[1.7]])
```

Out[43]:

```
array([1])
```

In [44]:

```
log_reg.predict_proba([[1.7]])
```

Out[44]:

```
array([[0.43834057, 0.56165943]])
```

## Exercise 3.1

In [45]:

```
multiclass_logreg_obj = LogisticRegression(multi_class="multinomial", solver="lbfgs", C=10)
# let us get the petal width. It is present in the 4th column of data
X = iris["data"][ :, [2,3]]
y = iris["target"]
```

In [46]:

```
print(X.shape)
print(y.shape)
```

```
(150, 2)
(150,)
```

In [47]:

```
df = pd.DataFrame(X, columns=[iris.feature_names[2], iris.feature_names[3]])
```

In [48]:

```
print(df.head())
```

	petal length (cm)	petal width (cm)
0	1.4	0.1
1	1.4	0.1
2	1.3	0.1
3	1.5	0.1
4	1.4	0.1

In [49]:

```
multiclass_logreg_obj.fit(X, y)
```

Out[49]:

```
LogisticRegression(C=10, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='multinomial',
    n_jobs=1, penalty='l2', random_state=None, solver='lbfgs',
    tol=0.0001, verbose=0, warm_start=False)
```

In [69]:

```
X_new = np.array([1, 0.1], ndmin=2)
prediction = multiclass_logreg_obj.predict(X_new)
pred_class = prediction[0]
pred_prob = multiclass_logreg_obj.predict_proba(X_new)
pred_prob_class = (pred_prob[0][pred_class])
print('Prediction {}'.format(prediction) + ', with a probability of: {}'.format(pred_prob_class))
```

Prediction [0], with a probability of: 0.9995796959204409